

VGGNet & ResNet final-report

김태훈¹

¹부산대학교 전기컴퓨터공학부.




Figure 1: Configuration of VGG-16

1 Introduction

VGGNet[7] and ResNet[1] are CNN models that improve the performance of previous CNN models using increasing depth and residual learning framework respectively. VGGNet won the first place on the ImageNet Large-Scale Visualization Challenge(ILSVRC) 2014, and ResNet won the first place on the ILSVRC 2015.

In this report, we use Pytorch to implement VGGNet and ResNet, and use CIFAR-10[4] datasets to train the CNN models for image classification, analyze and discuss the results.

2 VGGNet

2.1 Dataset

We use Images from the plane, car and bird categories in CIFAR-10 datasets. Each class consists of 6000 32×32 color images. There are 5000 training images and 1000 test images in Each class. therefore there are 15000 training images and 3000 test images in total.

2.2 Network architecture

We use Type-D configuration in the paper[7], commonly known as VGG-16. Type-D configuration is shown on 1.

Since we use 32×32 CIFAR-10 images, the last 3 FC(fully connected) layer are different from the VGGNet in the paper[7] that using Imagenet dataset resized into 224×224 . 32×32 images change to $1 \times 1 \times 512$ feature maps after VGG16 convolutional layers. furthermore, We use 3 categories images. Therefore, Our VGG16 model have 1 FC layer that classify input to 3 categories¹

All hidden layers are equipped with the ReLU[5], and batch normalization[2] is incorporated after every convolution, which helps in stabilizing and accelerating the training process, and which is not used in the paper[7]. Thus, after each convolution, batch normalization and ReLU are applied in order. Also, there are no dropout layers for simplicity. The number of parameters is about 14.7M.

2.3 Loss function

We use Cross-entropy loss between outputs and targets as loss function. Since nn.CrossEntropyLoss function takes logits before softmax as outputs and scaler-index as targets, We don't apply a softmax function after the output layer.

2.4 Training

The training is carried out by optimising the multiple logistic regression, using mini-match stochastic gradient descent optimizer with momentum. The optimizing method is back-propagation[6]. The batch size was set to 128, momentum to 0.9, learning rate to 10^{-2} , and weight-decay to 5×10^{-4} . Additionally, we trained the models for 20 epochs without learning rate scheduling. For simplicity, weight was initialized with default Pytorch weight initialization method.

VGG-16 with (mini-) CIFAR-10 config		
layer name	output size	
conv1_x	32x32x64	(3x3, 64) x2
pool1	16x16x64	maxpool
conv2_x	16x16x128	(3x3, 128) x2
pool2	8x8x128	maxpool
conv3_x	8x8x256	(3x3, 256) x3
pool3	4x4x256	maxpool
conv4_x	4x4x512	(3x3, 512) x3
pool4	2x2x512	maxpool
conv5_x	2x2x512	(3x3, 512) x3
pool5	1x1x512	maxpool
fc-3	3	(512, 3)

Table 1: Configuration of VGG-16 with CIFAR-10 config

The training images size was fixed to 32×32 . The training images was preprocessed with random cropping to 32×32 (The images was padded 4 for all direction before cropped), random horizontal flipping, changing images array to Pytorch tensors(which changes $Height \times Width \times Channel$ to $Channel \times Height \times Width$ and scale RGB values to 0 to 1), and normalizing pixels with mean and standard of RGB values of training images.

2.5 Testing

We tested the models for every epochs with test images. Test images was preprocessed with changing images array to Pytorch tensors and normalizing pixels with mean and standard of RGB values of training images. Random cropping and horizontal flipping was not applied to the test images because of consistency of testing. Testing method is classification accuracy(i.e., the precentage of correct prediction.

2.6 Implementation Details

2.6.1 Configuration array

The Configuration array(declared a 'cfg' on code) defines the architecture configuration of the VGG16 model. It specified the number of output channels for each convolutional layer and the placement of max pooling layers('MP').

2.6.2 Model Class

The Model Class, declared a 'VGG' on code, inherits from nn.Module. Attributes of the class defines the convolutional layers and fully connected layers. The convolutional layers declared a attribute 'VGG16': The make_layers method constructs the convolutional part of the network based on the cfg configuration. For each integer value in cfg, nn.Conv2d with output channels is added, followed by nn.BatchNorm2d and nn.ReLU. When a 'MP' is encountered, nn.MaxPool2d is added.

Fully Connected Layers declared a attribute 'classifier': there is one FC layer with 3 channels for classification tasks.

Forward Pass method consists of 3 steps: Convolutional Processing, Flattening, and Classification. Convolutional processing is that Input data x is passed through the convolutional layers. Flattening is that the output feature maps are reshaped into a 1D tensor using out.view(out.size(0), -1) for the fully connected layers. Classification is that the flattened tensor is passed the classifier to obtain the final logit values.

2.7 Results

After the final epochs, Training loss and Testing accuracy was about 0.0004 and 91.2 each³. Training loss and testing accuracy on every epoch show on fig²

3 ResNet

3.1 Dataset

We use the dataset that used for VGGNet.

3.2 Network Architecture

We use 50-layer ResNet(i.e., ResNet50) with bottleneck blocks. The configuration of ResNet50 is different from the paper[1], first is the initial convolutional layer changed from $7 \times 7, 64, stride = 2$ to $3 \times 3, 64, stride = 1$, and second is the initial max-pooling layer is removed because the size of CIFAR-10[4] images is too small.

The configuration of ResNet50 for CIFAR-10[4] is 4, and the residual connection is 5. All hidden layers(except the convolution layer for expanding channels) are equipped with the ReLU[5], and batch normal-ization[2] is incorporated after every convolution. Strided convolution for down-sampling instead of max-pooling layer. if Once down-sampled, a 1×1 convolution with stride 2 is applied to residual input for matching dimension of input and output. if down-sampling is not applied but the number of channels of input and output are different(i.e., the first block of cfg[0] blocks in 4), convolution with stride 1 is applied to residual input for expanding channels of input. Also, there are no dropout for simplicity. The number of parameters is about 23.5M.

3.3 Loss Function

The loss function is same with VGGNet above^{2.3}. In other words, Cross-entropy loss between outputs and ground-truths.

3.4 training

The training method is same with VGGNet above^{2.4}, except the number of epochs is 15.

3.5 implementation Details

3.5.1 configuration array

The configuration array('cfg' variable in code) defines the number of residual blocks in each of the four main residual net layers of the ResNet50. [3,4,6,3] indicates that Layer1-Layer4 has 3, 4, 6, 3 residual blocks each.

3.5.2 ResNet block class

ResNet block class defines a single residual block. The arguments of initialization of resnet block class is 'in-c', 'intra-c', 'out-c', 'down-sample', 'expand'. Each block consists of three convolution layers:

First Convolution conv1 consists of 1×1 convolution, Batch Normalization, and ReLU. 1×1 convolution reduces number of channels from 'in-c' to 'intra-c'. if 'down-sample' arguments is true, then stride of convolution is 2 for downsampling.

Second Convolution conv2 consists of 3×3 convolution, and others are same with conv1. 3×3 convolution layer maintain the number of channels(intra-c), and padding='same' to preserve spatial dimensions.

Third Convolution conv3 consists of 1×1 convolution, and Batch Normalization. 1×1 convolution expands the number of channels from intra-c to out-c.

If down-sample or expand attribute is true, the input undergoes a 1×1 convolution to match the dimensions of the output from the convolution layers. if down-sample is true, the stride of 1×1 convolution is 2(if not, then stride is 1). The output of the convolution layers is added to the input, and a ReLU activation is applied.

VGG-16 with (mini-) CIFAR-10 config		
layer name	output size	
conv1_x	32x32x64	(3x3, 64) stride 1
conv2_x	32x32x256	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	16x16x512	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$
conv4_x	8x8x1024	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$
conv5_x	4x4x2048	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
avgpool	2048	
fc1	3	

Table 2: Configuration of ResNet50 with CIFAR-10 config

	VGG16	ResNet50
Train Loss	0.0004	0.0586
Test Accuracy(%)	91.2	85.8333

Table 3: Train loss and test accuracy of VGG16 and ResNet50. the values was rounded off to the the fourth decimal place.

this residual connection enables the network to learn identity mapping and thus allowing the construction of deeper networks without vanishing or exploding gradient.

3.5.3 ResNet Model Class

The ResNet class defines the entire ResNet50 architecture. The class consists of initial block, residual layers, adaptive average pooling and classifier. initial block consists of 3×3 convolution that number of output channels is 64, Batch Normalization and ReLU activation. Residual layers are like ⁴²

The forward of this model has 4 part: Initial convolutional processing, passing residual layers, global average pooling, flattening, and classification.

Initial convolutional processing is that the input tensor passes through the initial block. passing residual layers is that the output from initial block sequentially passes through first to fourth residual block. global average pooling is that the output from the last residual layer undergoes adaptive average pooling, which changes feature map dimension from $Height \times Width \times Channels$ to $1 \times 1 \times Channel$. Finally, flattening and classification is that changes feature map to 1D tensor, and passed through classifier, which consists of a linear layers that maps the features to logit values.

3.6 Results

After the final epochs, Training loss and Testing accuracy was about 0.0586 and 85.8333 each³. Training loss and testing accuracy on every epoch show on fig³

4 Discussion

For both models, The training loss graph is fluctuating, which might suggest issues with the large learning rate. Applying the Adam optimizer[3] or setting learning rate to be smaller may resolve the issue. The test accuracy of both models are higher than 85%, which shows that deeper layer improve the performance. Deeper layers have ability to represent complex function, and hierarchical learning. VGG16 is better than ResNet50 in training loss and testing accuracy. That may be because ResNet has more parameters and deeper layers than VGGNet, making 15,000 training data insufficient to train all of its parameters. If there are more training data(and more categories), ResNet may be better than VGGNet because ResNet50 is better than VGG16 in the ResNet paper[1], which trained the models with ImageNet.

- [1] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition, 2015. URL <https://arxiv.org/abs/1512.03385>.
- [2] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift, 2015. URL <https://arxiv.org/abs/1502.03167>.
- [3] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization, 2017. URL <https://arxiv.org/abs/1412.6980>.
- [4] Alex Krizhevsky. Learning multiple layers of features from tiny images. *University of Toronto*, 05 2012.
- [5] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In F. Pereira, C.J. Burges, L. Bottou, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems*, volume 25. Curran Associates, Inc., 2012. URL https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf.
- [6] Y. LeCun, B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, and L. D. Jackel. Backpropagation applied to handwritten zip code recognition. *Neural Computation*, 1(4):541–551, 1989. doi: 10.1162/neco.1989.1.4.541.
- [7] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition, 2015. URL <https://arxiv.org/abs/1409.1556>.